

Data Shapley : Equitable Valuation of Data for Machine Learning

2023.06.20 (Tue.)

Superb AI Machine Learning Team

Presenter : Kyeongryeol, Go

Overview

- Data generated by individuals is a key component of the market place
 - E.g. health care, advertising
- 3 ingredients in case of supervised machine learning
 - A fixed training dataset : $D = \{(x_i, y_i)\}_{i=1}^n$
 - A learning algorithm : \mathcal{A} , which solves $\theta^* = \arg \min_{\theta} \sum l(f(x_i; \theta), y_i)$
 - A performance metric : V , test performance of f
 - Mean Squared Error in regression
 - Accuracy in classification

1. What is an equitable measure of the value of each (x_i, y_i) to \mathcal{A} w.r.t. V
2. How do we efficiently compute this data value in practical settings

Equitable data valuation

- Notation
 - $\phi_i(V)$: value of i -th train datum in terms of V
 - $V(S)$: test performance of f trained on $S \subseteq D$
c.f. $V(S = \Phi)$ indicates test performance of randomly initialized classifier
- What properties make ϕ_i “equitable”
 1. Null player : $\forall S \subseteq D - \{i\}, V(S \cup \{i\}) = V(S)$, then $\phi_i = 0$
 2. Symmetry : $\forall S \subseteq D - \{i, j\}, V(S \cup \{i\}) = V(S \cup \{j\})$, then $\phi_i = \phi_j$
 3. Linearity : $\forall \alpha_1, \alpha_2 \in \mathbb{R}, \phi_i(\alpha_1 V_1 + \alpha_2 V_2) = \alpha_1 \phi_i(V_1) + \alpha_2 \phi_i(V_2)$
 4. Efficiency : $V(D) = \sum_i \phi_i(V)$

Shapely is a unique measure satisfying the above four properties

Shapely formulation

- Formulation

$$(\text{Shapely}) \ \phi_i := \frac{1}{n} \sum_{S \subseteq D - \{i\}} \binom{n-1}{|S|}^{-1} (V(S \cup \{i\}) - V(S))$$

- Sum of **marginal gains** divided by **# of subsets with cardinality $|S|$**
 - assign uniform weights to different subset size $|S|$
 - note that such normalization is derived from the efficiency property

1. Traversing all the possible $S \subseteq D - \{i\}$ is computationally infeasible for large dataset
2. Every computation of V requires a separate model training

Approximating ϕ

- Monte-Carlo method

$$\phi_i := \frac{1}{n} \sum_{S \subseteq D - \{i\}} \binom{n-1}{|S|}^{-1} (V(S \cup \{i\}) - V(S))$$

$$= \frac{1}{n} \sum_{j=0}^{n-1} \binom{n-1}{j}^{-1} \sum_{\substack{S \subseteq D - \{i\} \\ |S|=j}} (V(S \cup \{i\}) - V(S))$$

$$= \mathbb{E}_{j \sim U(0, n-1)} \mathbb{E}_{\substack{S \subseteq D - \{i\} \\ |S|=j}} [V(S \cup \{i\}) - V(S)] = \mathbb{E}_{\pi \sim \Pi} [V(S_{\pi}^i \cup \{i\}) - V(S_{\pi}^i)]$$

- π : random permutation of indices $[n]$ -----
 - S_{π}^i : subset of indices coming before datum i -----
- > If $\pi = [2, 1, 5, 3, 4]$, then $S_{\pi}^5 = \{2, 1\}$

1. Sample a random permutation
2. Scan the permutation from the first to the last element and calculate the marginal gain (truncate the calculation of marginal gain whenever $|V(D) - V(S_{\pi}^i)| < tolerance$)
3. Repeat 1 and 2, then obtain final estimation by average

$$t=1, \pi = [2, 1, 5, 3, 4]$$

$$\phi_2^1 = V(\{2\})$$

$$\phi_1^1 = V(\{2, 1\}) - V(\{2\})$$

...

$$\phi_4^1 = V(\{2, 1, 5, 3, 4\}) - V(\{2, 1, 5, 3\})$$

...

$$t=T, \pi = [5, 3, 1, 2, 4]$$

$$\phi_5^T = V(\{5\})$$

$$\phi_3^T = V(\{5, 3\}) - V(\{5\})$$

...

$$\phi_4^T = V(\{5, 3, 1, 2, 4\}) - V(\{5, 3, 1, 2\})$$

$$\phi_i = \frac{1}{T} \sum_{t=1}^T \phi_i^t \quad \forall i \in [5]$$

Approximating V

1. Train the model with only one epoch and with mini-batch size of 1 → Algorithm 2

- marginal gain can be computed between iterations
- require HPO to find the one resulting best performance by one epoch (high learning rate)

2. Estimate value by group of data points → Algorithm 1

e.g. group the patients into discrete bins based on age, gender, ethnicity, etc.

Algorithm 1 Truncated Monte Carlo Shapley

Input: Train data $D = \{1, \dots, n\}$, learning algorithm \mathcal{A} , performance score V

Output: Shapley value of training points: ϕ_1, \dots, ϕ_n

Initialize $\phi_i = 0$ for $i = 1, \dots, n$ and $t = 0$

while Convergence criteria not met **do**

$t \leftarrow t + 1$

π^t : Random permutation of train data points

$v_0^t \leftarrow V(\emptyset, \mathcal{A})$

for $j \in \{1, \dots, n\}$ **do**

if $|V(D) - v_{j-1}^t| < \text{Performance Tolerance}$ **then**

$v_j^t = v_{j-1}^t$ ↳ 0 value for the rest of j

else

$v_j^t \leftarrow V(\{\pi^t[1], \dots, \pi^t[j]\}, \mathcal{A})$

end if

$\phi_{\pi^t[j]} \leftarrow \frac{t-1}{t} \phi_{\pi^{t-1}[j]} + \frac{1}{t} (v_j^t - v_{j-1}^t)$

end for

end for

Algorithm 2 Gradient Shapley

Input: Parametrized and differentiable loss function $\mathcal{L}(\cdot; \theta)$, train data $D = \{1, \dots, n\}$, performance score function $V(\theta)$

Output: Shapley value of training points: ϕ_1, \dots, ϕ_n

Initialize $\phi_i = 0$ for $i = 1, \dots, n$ and $t = 0$

while Convergence criteria not met **do**

$t \leftarrow t + 1$

π^t : Random permutation of train data points

$\theta_0^t \leftarrow \text{Random parameters}$

$v_0^t \leftarrow V(\theta_0^t)$

for $j \in \{1, \dots, n\}$ **do**

$\theta_j^t \leftarrow \theta_{j-1}^t - \alpha \nabla_{\theta} \mathcal{L}(\pi^t[j]; \theta_{j-1})$

$v_j^t \leftarrow V(\theta_j^t)$

$\phi_{\pi^t[j]} \leftarrow \frac{t-1}{t} \phi_{\pi^{t-1}[j]} + \frac{1}{t} (v_j^t - v_{j-1}^t)$

end for


end for

Convergence criteria
(for generating random permutation)


$$\Rightarrow \frac{1}{n} \sum_{i=1}^n \frac{|\phi_i^t - \phi_i^{t-100}|}{|\phi_i^t|} < 0.05$$

< From lowest valued data >

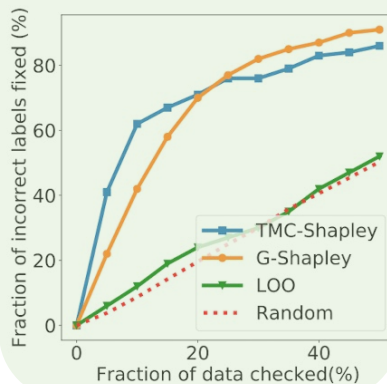
Spam Classification
Naïve Bayes Classifier
20% mislabeled






Flower Classification
Multinomial Logistic Regression
10% mislabeled





T-Shirt/Top vs Shirt Classification
ConvNet Classifier
10% mislabeled



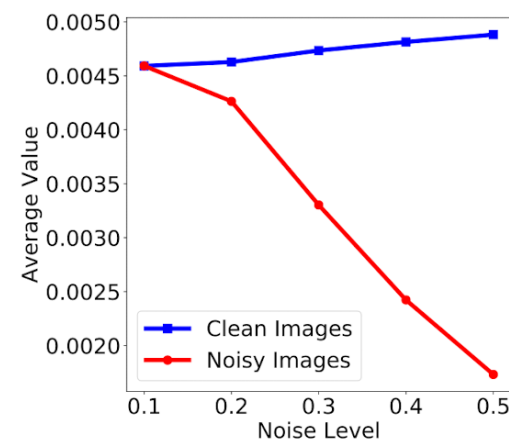
Flowers



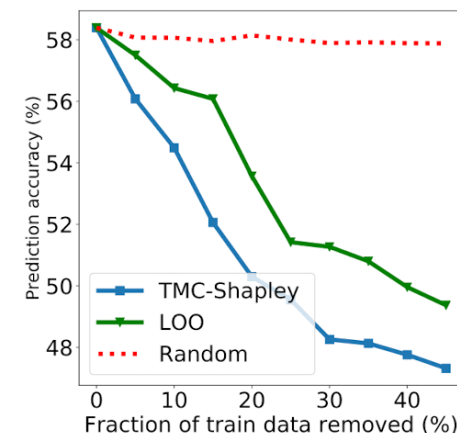
Fashion MNIST



5 images with the lowest value : all mis-labeled



< From largest valued group >



2 classes
3000 data points
5000 iterations

5 classes
1000 images
2000 iterations

2 classes
1000 images
2000 iterations

Use representation from Inception-V3 model

Use one convolution + 2 feed-forward layers

1. Mis-labeled data has low values
2. Discrepancy b/t clean and noisy gets larger as noise level increases
3. Deletion of groups of high values leads to huge performance drop

Use Inception-V3 model unfreezing the last layer

Use a gradient boosting classifier

2 classes
100 images

2 classes
60000 data points
146 groups

Distributional Shapely

- Data Shapely assumed a fixed dataset D
 - no guarantee on consistency between the different, but similar datasets
 - it may be very sensitive to the exact choice of D and not transferrable across datasets
- Idea : Suggest an unbiased estimator for the size of D

$$v_i = \mathbb{E}_{B \subseteq D} [\phi_i] = \mathbb{E}_{B \subseteq D} \mathbb{E}_{j \sim U(0, |B| - 1)} \mathbb{E}_{\substack{S \subseteq B - \{i\} \\ |S|=j}} [V(S \cup \{i\}) - V(S)]$$

- more stable under perturbations to inputs as well as the underlying data distribution
- **Thm.** under Lipschitz Stable performance metric V , the following holds
(c.f. $|V(S \cup \{i\}) - V(S \cup \{j\})| \leq \beta(|S|) \cdot d(i, j)$)
 1. Similar distributions yield similar value functions $\Rightarrow |v_i(D_s) - v_i(D_t)| \leq C \cdot W(D_s, D_t)$
 2. Similar points receive similar values $\Rightarrow |v_i - v_j| \leq C \cdot d(i, j)$

Distributional Shapely

- Distributional Shapely assigns relatively larger weights to small sized subsets

(note : (i) $S \subseteq B - \{i\}$, $B \subseteq D$, (ii) Data shapely : $B = D$)

Data shapely →

weight	$ S =0$	$ S =1$...	$ S =n-2$	$ S =n-1$	sum
$ B =n$	1/n	1/n	...	1/n	1/n	1
$ B =n-1$	1/(n-1)	1/(n-1)	...	1/(n-1)	0	1
...
$ B =1$	1/2	1/2	...	0	0	1
$ B =0$	1	0	...	0	0	1

larger
smaller

- no more equal weights to subset size → no more satisfy the efficiency property := semi-values
- It is controversial whether it is beneficial to satisfy the efficiency in machine learning
 - ∴ the value itself doesn't matter, but the rank matters

Beta Shapely

- Data Shapely assumed a uniform weight over the subset size $|S|$
 - Thm.** When $|S|$ is large enough, the marginal gain becomes negligible (signal-to-noise ratio($:= \mu/\sigma$) increases as the cardinality increases)
(more likely to be perturbed by noise)



- General representation of semi-value

$$\psi_i = \frac{1}{n} \sum_{j=0}^{n-1} w(j) \sum_{\substack{S \subseteq D - \{i\} \\ |S|=j}} V(S \cup \{i\}) - V(S) \quad \text{s.t.} \quad \sum_{j=0}^{n-1} \binom{n-1}{j} w(j) = n$$

$$\text{(c.f. } \phi_i = \frac{1}{n} \sum_{j=0}^{n-1} \binom{n-1}{j}^{-1} \sum_{\substack{S \subseteq D - \{i\} \\ |S|=j}} V(S \cup \{i\}) - V(S) \text{)}$$

- Idea : Assign large weights to small cardinality

- $w_{\alpha,\beta}(j) = n \frac{\text{Beta}(j+\beta, n-j-1+\alpha)}{\text{Beta}(\alpha,\beta)} \Rightarrow (\alpha,\beta) = (1,1)$ suffices to Data Shapely
- $\alpha \geq \beta = 1$ assigns large weights on the small cardinality and remove noise from large cardinality



Data Banzhaf

- Stochastic training methods (e.g. SGD) make the performance metric unreliable
 - noise is shown to be substantial to make different runs produce inconsistent value rankings
- Idea 1 : Suggest to use Banzhaf value as a robust data value notion
 - $\psi_i = \frac{1}{n} \sum_{j=0}^{n-1} \frac{n}{2^{n-1}} \sum_{\substack{S \subseteq D - \{i\} \\ |S|=j}} V(S \cup \{i\}) - V(S) = \mathbb{E}_{S \subseteq D - \{i\}} [V(S \cup \{i\}) - V(S)]$
 - Note that it is also the semi-value s.t. $\sum_{j=0}^{n-1} \binom{n-1}{j} \frac{n}{2^{n-1}} = n$
 - $\frac{n}{2^{n-1}}$ is not a function of cardinality, which leads large weights to subsets of large cardinality
 - **Thm.** Banzhaf value achieves the largest safety margin among the semi-values
 - Safety margin : the largest tolerable perturbation on V that keeps the value order unchanged
- Idea 2 : Maximum Sample Reuse (MSR) Monte Carlo
 - Sample a number of subsets with various cardinalities $\mathcal{S} = \{S_1, \dots, S_m\} = \mathcal{S}_{\ni i} \cup \mathcal{S}_{\not\ni i}$
 - $\psi_i = \frac{1}{|\mathcal{S}_{\ni i}|} \sum_{S \in \mathcal{S}_{\ni i}} V(S) - \frac{1}{|\mathcal{S}_{\not\ni i}|} \sum_{S \in \mathcal{S}_{\not\ni i}} V(S) \Rightarrow$ enable faster estimation

E.O.D