

Never Give Up Learning Directed Exploration Strategies


Accepted in ICLR 2020

Kyeong Ryeol, Go
M.S. Candidate of OSI Lab

Overview


- Motivation

- In sparse reward setting, exploration strategy is crucial




$$100 * (DQNscore - randomplayscore) / (humanscore - randomplayscore)$$

Overview



Break out



Montezuma's
revenge



Pitfalls!

Overview

- Intrinsic reward
 - Quantifies the novelty of the experience to respond to agent's curiosity
- Previous work
 - Visitation counts : ignore the downstream learning opportunity
 - Prediction error : expensive, error prone, generalization issue
- Main idea
 - Designing the intrinsic rewards to encourage the agents to visit diverse states **within** and **across** episodes
 - Learn a range of policies with varying the exploration and exploitation trade-offs

Contents

1. Never Give Up Intrinsic reward
2. Model architecture & Loss function
3. Experiment
4. Conclusion
5. Appendix

Never Give Up Intrinsic reward

NGU Intrinsic Reward

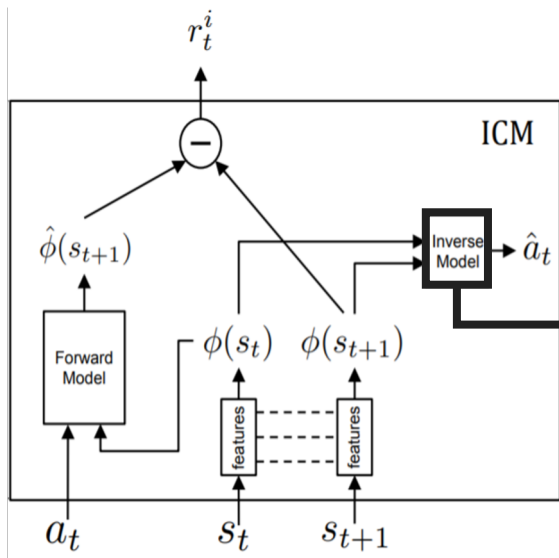
- Reward at time t : $r_t = r_t^e + \beta r_t^i$
 - Linear combination of extrinsic reward and intrinsic reward
 - β is the balancing parameter
- 3 Properties of intrinsic reward
 - The notion of state ignores aspects of an environment that are not influenced by an agent.
(Embedding function)
 - It rapidly discourages revisiting the same state within the same episode.
(Episodic Novelty Module)
 - It slowly discourages visiting to states where visited many times across episodes.
(Life-long Novelty Module)

NGU Intrinsic Reward - Embedding function


- Embedding function f

- The notion of state ignores aspects of an environment that are not influenced by an agent.
 - Input : two consecutive observations
 - Output : In-between action

→ Robust to inherent stochasticity



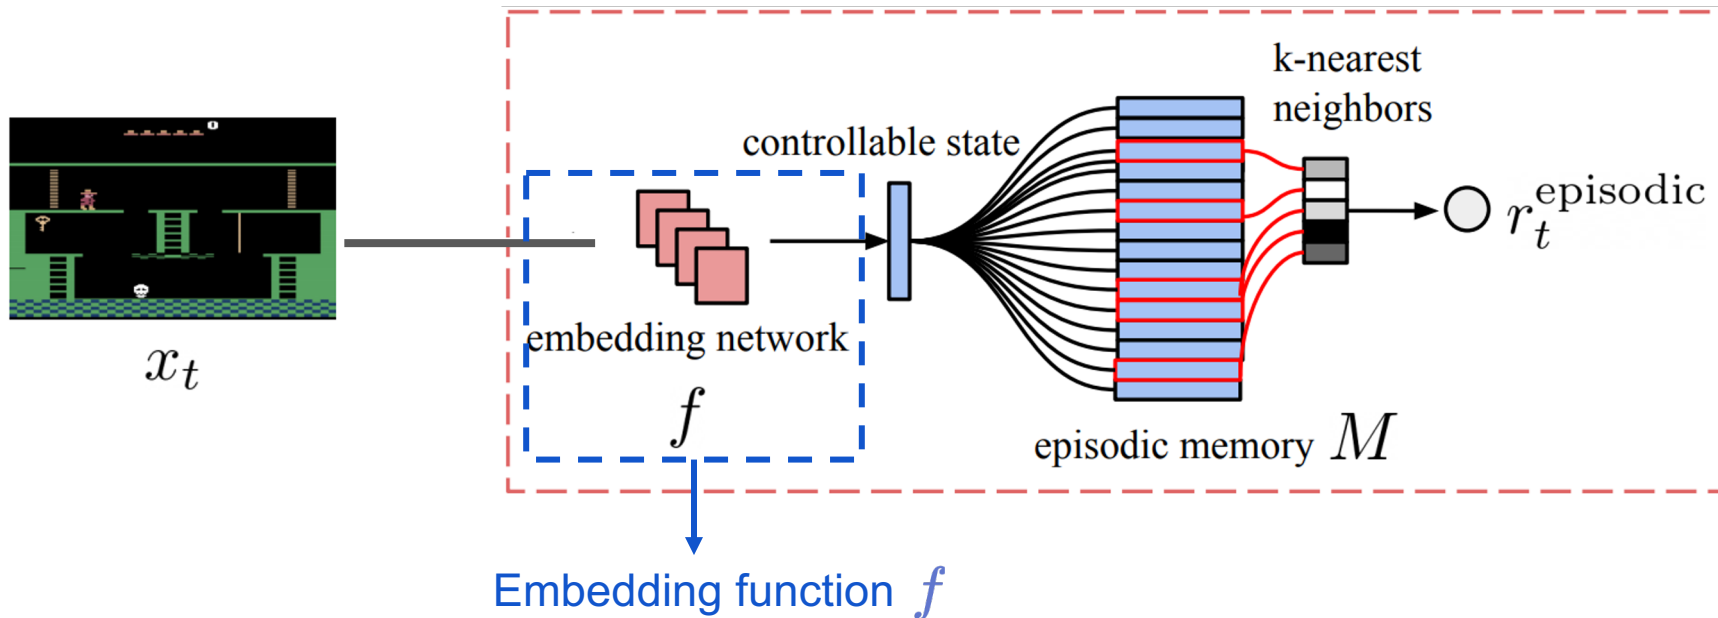
Intrinsic Curiosity Module (ICM)



NGU Intrinsic reward - Episodic novelty module

- Episodic novelty module r_t^{episodic}

- It rapidly discourages revisiting the same state within the same episode.



NGU Intrinsic reward - Episodic novelty module

- Episodic novelty module $r_t^{episodic}$

- Episodic memory M : At every step, agent computes an episodic intrinsic reward $r_t^{episodic}$ and appends the controllable state corresponding to the current observation to the memory M

$$\boxed{r_t^{episodic} = \frac{1}{\sqrt{n(f(x_t))}}} \approx \frac{1}{\sqrt{\sum_{f_i \in N_k} K(f(x_t), f_i) + c}} \quad K(x, y) = \frac{\epsilon}{\frac{d^2(x, y)}{d_m^2} + \epsilon}$$


From UCB

- embedding function $f : \mathcal{O} \rightarrow \mathbb{R}^p$: mapping the current observation to a p-dimensional vector corresponding to its controllable state
- episodic novelty : Promotes the agent to visit as many different states as possible within a single episode

NGU Intrinsic Reward - Life-long novelty module

- Life-long novelty module α_t

- It slowly discourages visiting to states where visited many times across episodes.
- SOTA for Montezuma's revenge at the time



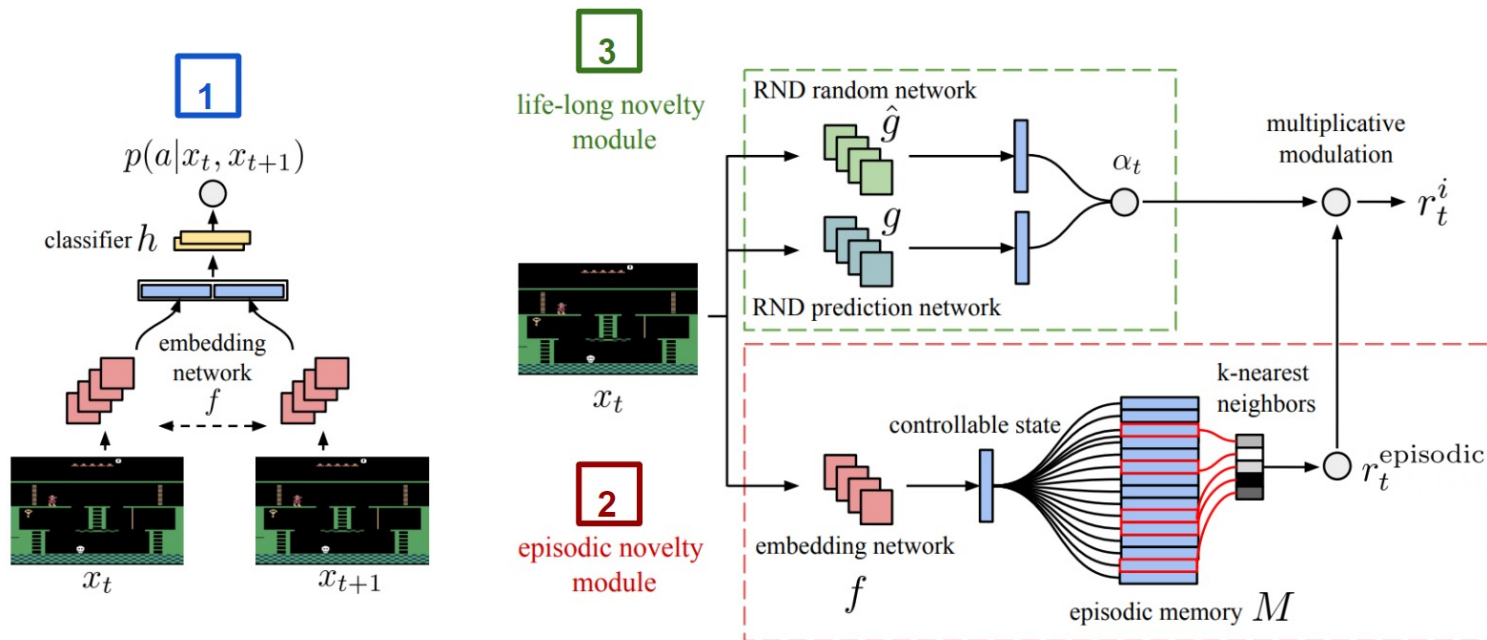
$$\text{err}(x_t) = \|\hat{g}(x_t; \theta) - g(x_t)\|^2$$

$$\alpha_t = 1 + \frac{\text{err}(x_t) - \mu_e}{\sigma_e}$$

Error would be small if similar experience to the current one would have been accumulated

NGU Intrinsic Reward

- Intrinsic reward $r_t^i = r_t^{\text{episodic}} \cdot \min\{\max\{\alpha_t, 1\}, L\}$
 - Episodic intrinsic reward is modulated by Life-long Intrinsic reward



Model architecture & Loss function

NGU - Proposed architecture


- Train the function approximator for $Q(x, a)$
 - Utilize UVFA to simultaneously consider a range of policies varying the weights on the intrinsic reward for overall reward design $\{\beta_i\}_{i=0}^{N-1} : r_t^{\beta_i} = r_t^e + \beta_i r_t^i$

※ Recurrent Replay Distributed QN (R2D2)

: DQN with LSTM module with distributed prioritized experience replay from Ape-X

※ Duel Deep Q-Network (Duel DQN)

: Devise two stream model whose outputs are $V(s)$ and $A(s, a)$



NGU - RL Loss function (1)

- Retrace(λ)
 - safe and efficient off-policy algorithm by using $c_k = \lambda \min\left(1, \frac{\pi(a_k, x_k)}{\mu(a_k, x_k)}\right)$
 - π : target policy, μ : behavior policy
- (Used along with Double q-learning and multi-step return to lower the variance)

Arbitrary behavior policy is fine

$$y_t^{\text{Retrace}(\lambda)} = E_{a_{t+1}, \dots, a_{t+s} \sim \mu} \left[\sum_{s=0}^k \gamma^s \left(\prod_{i=1}^s c_{t+i} \right) \delta_s + Q(x_t, a_t; \theta^-) \right]$$

$$\delta_s = r_{t+s} + \gamma E_{a_{t+s+1} \sim \pi} [Q(x_{t+s+1}, a_{t+s+1}; \theta^-)] - Q(x_t, a_t; \theta^-)$$

Select action increasingly greedy w.r.t. $Q(x, a; \theta)$

Evaluate target on $Q(x, a; \theta^-)$

NGU - RL Loss function (2)

- Transformed Retrace(λ)
 - Squash the scale of the action-value function

$$y_t^{T_Retrace(\lambda)} = E_{a_{t+1}, \dots, a_{t+s} \sim \mu} \left[h \left(\sum_{s=0}^k \gamma^s \left(\prod_{i=1}^s c_{t+i} \right) \delta_s^h + h^{-1}(Q(x_t, a_t; \theta^-)) \right) \right]$$

$$\delta_s^h = r_{t+s} + \gamma E_{a_{t+s+1} \sim \pi} [h^{-1}(Q(x_{t+s+1}, a_{t+s+1}; \theta^-))] - h^{-1}(Q(x_t, a_t; \theta^-))$$

$$h(x) = \text{sign}(x) \left(\sqrt{|x| + 1} - 1 \right) + \epsilon x,$$

$$h^{-1}(x) = \text{sign}(x) \left(\left(\frac{\sqrt{1 + 4\epsilon(|x| + 1 + \epsilon)} - 1}{2\epsilon} \right) - 1 \right)$$

$$L(\theta) = E_{(s,a,r,s') \in B} \left[\left(y^{T_Retrace(\lambda)} - Q(s, a; \theta) \right)^2 \right]$$


Experiment

Experiments

- Random Disco Maze
 - Things to check
 1. Effectiveness of the purely exploratory policy
 2. Effectiveness of the controllable state representation





Black – pathways, **Green** – agent
Random - wall



Experiments

- Ablation study on ATARI

- N \approx Number of beta values (Default = 32)
- CMR \approx Proportion of off-policy experiences (Default = 0)
- β \approx Magnitude of the maximum beta value (Default = 0.3)
- RND \approx Whether the life-long novelty module is used or not (Default = with)
- r^e \approx Whether the extrinsic reward is used or not (Default = with)



Justification

- UVFA for varying β
- RND for modulation
- Augmented reward

Experiments

- Comparing to several baselines

Hard Exploration Games	Algorithm	Gravitar	MR	Pitfall!	PrivateEye	Solaris	Venture
	Human	3.4k	4.8k	6.5k	69.6k	12.3k	1.2k
	Best baseline	15.7k	11.6k	0.0	11k	5.5k	2.0k
	RND	3.9k	10.1k	-3	8.7k	3.3k	1.9k
	R2D2+RND	15.6k±0.6k	10.4k±1.2k	-0.5±0.3	19.5k±3.5k	4.3k±0.6k	2.7k±0.0k
	R2D2(Retrace)	13.3k±0.6k	2.3k±0.4k	-3.5±1.2	32.5k±4.7k	6.0k±1.1k	2.0k±0.0k
	NGU(N=1)-RND	12.4k±0.8k	3.0k±0.0k	15.2k±9.4k	40.6k±0.0k	5.7k±1.8k	46.4±37.9
	NGU(N=1)	11.0k±0.7k	8.7k±1.2k	9.4k±2.2k	60.6k±16.3k	5.9k±1.6k	876.3±114.5
	NGU(N=32)	14.1k±0.5k	10.4k±1.6k	8.4k±4.5k	100.0k±0.4k	4.9k±0.3k	1.7k±0.1k
Dense Reward Games	Algorithm	Pong	QBert	Breakout	Space Invaders	Beam Rider	
	Human	14.6	13.4k	30.5	1.6k	16.9k	
	R2D2	21.0	408.8k	837.7	43.2k	188.2k	
	R2D2+RND	20.7±0.0	353.5k±41.0k	815.8±5.3	54.5k±2.8k	85.7k±9.0k	
	R2D2(Retrace)	20.9±0.0	415.6k±55.8k	838.3±7.0	35.0k±13.0k	111.1k±5.0k	
	NGU(N=1)-RND	-8.1±1.7	647.1k±50.5k	864.0±0.0	45.3k±4.9k	166.5k±8.6k	
	NGU(N=1)	-9.4±2.6	684.7k±8.8k	864.0±0.0	43.0k±3.9k	114.6k±2.3k	
	NGU(N=32)	19.6±0.1	465.8k±84.9k	532.8±16.5	44.6k±1.2k	68.7k±11.1k	

Conclusion

Conclusion

- Summary
 - Designing the intrinsic rewards to encourage the agents to visit diverse states **within** and **across** episodes
([Embedding function](#), [Episodic Novelty Module](#), [Life-long Novelty Module](#))
 - Learn a range of policies with varying the exploration and exploitation trade-offs
([UVFA framework](#))
- Extension
 - Agent 57 : Outperforming the Atari human benchmark
(<https://arxiv.org/abs/2003.13350>)
 - Utilizing the meta controller to choose which beta to be utilized for each episode

Thank you for your attention

Appendix

DQN variants

- Deep Q-Network (DQN)

- Utilize DNN as the function approximator for Q-value function
- Utilize the experience replay and target network to stabilize learning

$$L(\theta) = E_{(s,a,r,s') \in B} \left[\left(y^{DQN} - Q(s, a; \theta) \right)^2 \right] \text{ where } y^{DQN} = r + \gamma \max_{a'} Q(s', a'; \theta^-)$$

- Double Deep Q-Network (DDQN)

- Resolve the overestimation by selecting the action by the online network and estimating Q-value by the target network

$$y^{DDQN} = r + \gamma Q \left(s', \arg \max_a Q(s', a'; \theta) ; \theta^- \right)$$

- Dueling Deep Q-Network (Duel DQN)

- Two streams of network to compute $V(s)$ and $A(s, a)$ for computing $Q(s, a)$

$$Q(s, a) = V(s) + \left(A(s, a) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a') \right)$$

Sampling with Priority

- Prioritized Experience Replay
 - Impose priority on the experiences by the absolute TD error
 - Use importance sampling weight to correct the bias

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha} \text{ where } p_i = |\delta_i| + \epsilon \text{ or } p_i = \frac{1}{\text{rank}(|\delta_i|)}$$

$$\delta = r + \gamma Q\left(s', \arg \max_a Q(s', a'; \theta); \theta^-\right) - Q(s, a; \theta)$$

$$\theta \leftarrow \theta + \eta \cdot \frac{w_i}{\max_i w_i} \cdot \delta_i \cdot \nabla_{\theta} Q(s_i, a_i) \text{ where } i \sim P(i) \text{ and } w_i = \left(\frac{1}{N \cdot P(i)}\right)^\beta$$

(continued)

- Distributed Prioritized Experience Replay (Ape-X DQN)
 - Extend prioritized experience replay to the distributed setting
 - Actors : Select actions in the environment and store them in a buffer
 - Set the priority by the absolute TD error
 - Periodically synchronize the parameter of the learner
 - Learner : Sample experiences with priority and update the policy parameter
 - Update the priority again by the absolute TD error with the updated parameter
 - Use double q-learning and multi-step target

$$y_t^{Ape-X} = \sum_{k=0}^{n-1} \gamma^k r_{t+k} + \gamma^n Q\left(s_{t+n}, \arg \max_{a'} Q(s_{t+n}, a'; \theta); \theta^-\right)$$

Recurrent Replay Distributed DQN

- Recurrent Replay Distributed DQN (R2D2)

- Resolve the representation drift and recurrent state staleness

- Storing recurrent state in replay memory

- Allow a burn-in period by the portion of replay memory

(verifies its effectiveness by checking the Q-value discrepancy)

- Impose priority by a mixture of max and mean absolute n-step TD-error

$$p(i) = \eta \cdot \max_i \delta_i + (1 - \eta) \cdot \text{mean}_i |\delta_i|$$

- Modify Ape-X DQN target by rescaling

$$y_t^{R2D2} = \sum_{k=0}^{n-1} \gamma^k r_{t+k} + \gamma^n h^{-1} \left(Q \left(s_{t+n}, \arg \max_{a'} Q(s_{t+n}, a'; \theta); \theta^- \right) \right)$$

$$h(x) = \text{sign}(x) \left(\sqrt{|x| + 1} - 1 \right) + \epsilon x, \quad h^{-1}(x) = \text{sign}(x) \left(\left(\frac{\sqrt{1 + 4\epsilon(|x| + 1 + \epsilon - 1)}}{2\epsilon} \right) - 1 \right)$$